

(This slide intentionally  
left blank)

# So you want to write an interpreter?

Alex Gaynor, PyCon 2013

Friday, March 15, 13

Hi everyone. Thanks for coming. This talk is going to be about how to write your own interpreter. This talk is going to go pretty quickly, there's a lot of material to cover, but all the code is going to be online, so don't worry about copying every word down.



**rdio.com**

Friday, March 15, 13

First, thanks to my employer, rdio.com, we're a streaming internet music service, it's really awesome, we're hiring, you should come talk to me.

# Who am I?

- DSF/PSF
- I work on PyPy, CPython, Django
- I created Topaz

# What is an interpreter?

- It is a program. Just a bunch of code like any other.
- Its input is a program in some language.
- Its output is the result of running that program.
- Not magic!

Friday, March 15, 13

What is an interpreter? A lot of people consider interpreters and compilers to be some sort of black magic, a dark art, that you have to be indoctrinated into or something. They're not (or maybe this is the indoctrination). They're just programs.

# Why do you want to write an interpreter?

- Create your own language?
- Implement an existing language?
- Have fun!

Friday, March 15, 13

Why would you want to write one? Well to create a new language, to reimplement an existing language (I like doing this one a lot). Most of all, it can be a ton of fun, there's lots of really cool mini-puzzles that go into this.

# The pieces

- Lexer
- Parser
- AST
- Bytecode-compiler
- Bytecode interpreter
- Runtime

On your marks...  
Get set...  
Go!

Friday, March 15, 13

Basically for the rest of this talk we're going to build a VM for a tiny subset of Javascript. From the ground up, it's going to be really awesome.



# Our language

```
a = 3;  
if (a >= 2) {  
    print "a is big!";  
}  
else {  
    print a;  
}
```

Friday, March 15, 13

Our language contains just a few features, but it's really enough to do almost anything we want. The features we have are: numbers (floats) and strings, a few binary operators for comparisons and arithmetic, an if/else statement, and a print statement (that's not part of javascript, I made it up).

```
pip install rply
```

Friday, March 15, 13

RPLY is a toolkit we're going to use to help us construct our parser and lexer.

# Lexers!

Friday, March 15, 13

So the very first step in writing an interpreter is writing a lexer.

# What is a lexer?

Friday, March 15, 13

A lexer basically takes a program, and turns it into a list of tokens. What is a token a token is a symbol in your program, let's take a look at an example.

```
a = 3;
if (a >= 2) {
    print "a is big!";
}
else {
    print a;
}
```



```
[
Token("NAME", "a"),
Token("EQUAL", "="),
Token("SEMICOLON", ";"),
Token("IF", "if"),
Token("LPAREN", "("),
Token("NAME", "a"),
Token("GREATER_EQUAL", ">="),
Token("NUMBER", "2"),
Token("LBRACE", "{"),
Token("PRINT", "print"),
Token("STRING", "a is big!"),
Token("SEMICOLON", ";"),
Token("RBRACE", "}"),
Token("ELSE", "else"),
Token("LBRACE", "{"),
Token("PRINT", "print"),
Token("NAME", "a"),
Token("SEMICOLON", ";"),
Token("RBRACE", "}"),

```

```
]
```

So that program, becomes that list of tokens. A token is basically two things, 1) it has a name, and b) it has the text it matches, so an “LBRACE” token always has the same text, but a NAME or NUMBER has different text. Also, notice that there’s no whitespace, it’s not significant so we just ignore it.

WHENEVER I LEARN A NEW SKILL I CONCOCT ELABORATE FANTASY SCENARIOS WHERE IT LETS ME SAVE THE DAY.

OH NO! THE KILLER MUST HAVE FOLLOWED HER ON VACATION!



BUT TO FIND THEM WE'D HAVE TO SEARCH THROUGH 200 MB OF EMAILS LOOKING FOR SOMETHING FORMATTED LIKE AN ADDRESS!



IT'S HOPELESS!

EVERYBODY STAND BACK.



I KNOW REGULAR EXPRESSIONS.



Friday, March 15, 13

How do we implement lexers? Regular expressions! Each token type is just a specific regular expression. Really simple.

```
from rply import LexerGenerator
lg = LexerGenerator()
lg.ignore(r"\s+")
lg.add("IF", r"if")
lg.add("ELSE", r"else")
lg.add("PRINT", r"print")
lg.add("LPAREN", r"\(")
lg.add("RPAREN", r"\)")
lg.add("LBRACE", r"\{")
lg.add("RBRACE", r"\}")
lg.add("EQUAL", r"=")
lg.add("GREATER_EQUAL", r">=")
lg.add("SEMICOLON", r";")
lg.add("NUMBER", r"\d+")
lg.add("NAME", r"[a-zA-z_][a-zA-Z0-9_]*")
lexer = lg.build()
```

Friday, March 15, 13

So this is basically how lexers in RPLY work. You create a generator, you add the rules you want, with a name and a regexp, and then you create

```
>>> from minijs.lexer import lexer
>>> stream = lexer.lex("my_var = 23")
>>> stream
<rply.lexer.LexerStream object at 0x10bfb2090>
>>> stream.next()
Token('NAME', 'my_var')
>>> stream.next()
Token('EQUAL', '=')
>>> stream.next()
Token('NUMBER', '23')
>>> stream.next()
>>>
```



# Parsers

Friday, March 15, 13


So in general, each stage of our interpreter is going to take the output of the previous stage and do something with it. So our first stage took the source code and gave us a list of tokens. Now we need to take a list of tokens and do something with it.

[

```

Token( "NAME" , "a" ),
Token( "EQUAL" , "=" ),
Token( "SEMICOLON" , ";" ),
Token( "IF" , "if" ),
Token( "LPAREN" , "(" ),
Token( "NAME" , "a" ),
Token( "GREATER_EQUAL" , ">=" ),
Token( "NUMBER" , "2" ),
Token( "LBRACE" , "{" ),
Token( "PRINT" , "print" ),
Token( "STRING" , "a is big!" ),
Token( "SEMICOLON" , ";" ),
Token( "RBRACE" , "}" ),
Token( "ELSE" , "else" ),
Token( "LBRACE" , "{" ),
Token( "PRINT" , "print" ),
Token( "NAME" , "a" ),
Token( "SEMICOLON" , ";" ),
Token( "RBRACE" , "}" ),

```



```

Block([
  Assignment( "a" , Number( 3 ) ),
  If(
    Comparison( ">=" , Name( "a" ) , Number( 2 ) ),
    Block([
      Print( String( "a is big!" ) ),
    ]),
    Block([
      Print( Name( "a" ) ),
    ])
  )
])

```

]

Friday, March 15, 13

So basically we take our list of tokens, and turn it into a tree which represents what's going on. A Block is just a group of operations. Assignment is assignment. Name is a variable load. etc. All the nodes are pretty simple. This is called an abstract syntax tree.

```

class Node(object):
    def __eq__(self, other):
        if not isinstance(other, Node):
            return NotImplemented
        return (type(self) is type(other) and
                self.__dict__ == other.__dict__)
    def __ne__(self, other):
        return not (self == other)
class Block(Node):
    def __init__(self, statements):
        self.statements = statements
class Statement(Node):
    def __init__(self, expr):
        self.expr = expr
class Number(Node):
    def __init__(self, value):
        self.value = value

```

Friday, March 15, 13

So this is our initial set of AST classes, sorry about the very un-pep8 whitespace. Basically each one of these is a wrapper for some other stuff. We'll see why we have Statement later once we get to compilation. Now we need the parser that actually parses.

```
from rply import ParserGenerator
from minijs import ast
pg = ParserGenerator([ "NUMBER", "SEMICOLON" ], cache_id="minijs")
parser = pg.build()

@pg.production("main : statements")
def main(s):
    return s[0]
```

Friday, March 15, 13

So basically the parser gets a set of rules, and an action to perform when there's a match for each rule. I consider this to be by far the most confusing and hardest part of writing an interpreter, so if it's confusing... welcome to the club.

```
@pg.production("statements : statements statement")
def statements(s):
    return ast.Block(s[0].getastlist() + [s[1]])

@pg.production("statements : statement")
def statements_statement(s):
    return ast.Block([s[0]])

@pg.production("statement : expr SEMICOLON")
def statement_expr(s):
    return ast.Statement(s[0])

@pg.production("expr : NUMBER")
def expr_number(s):
    return ast.Number(float(s[0].getstr()))
```

Friday, March 15, 13

If you look you can see how this builds up. A group of statements is either one statement, or a list of statements. A statement is an expression with a semicolon, an expression is a number. And of course we'd want to expand this. For example addition is also an expression. There's tons of things we don't have here.

```
@pg.production("expr : expr PLUS expr")
def expr_binop(s):
    return ast.BinOp(s[1].getstr(), s[0], s[2])
```

Friday, March 15, 13

So this rule is a fun one. `expr + expr`. What can this parse? It can handle `1 + 1 + 1`, or anything else like that because the left and right hand sides *are* `exprs`, but the result is also an `expr`.

# Compiler

Friday, March 15, 13

So the next step is taking the AST we just produced, and turning it into something for an interpreter.

# Bytecode

- A sequence of instructions for the machine
- Sometimes have arguments.
- Stack based.



1 + 21;

LOAD\_CONST 0

LOAD\_CONST 1


BINARY\_ADD

POP\_TOP

Friday, March 15, 13

So how does this work, we've got a few instructions here. LOAD\_CONST, BINARY\_ADD, POP\_TOP. How does this work? We see the number 1. That becomes a LOAD\_CONST, what's that 0? The 0 means refer to the 0th constant, and we'll keep an array of them somewhere else. Then we load up const number 1 (the second const). We now have 2 items on our stack. Then we add them together, that pops 1 and 21 off the stack, adds them, and pushes the result.

1 + 21;




LOAD\_CONST 0

LOAD\_CONST 1


BINARY\_ADD

POP\_TOP



1 + 21;

→ LOAD\_CONST 0  
LOAD\_CONST 1  
BINARY\_ADD  
POP\_TOP




1 + 21;

LOAD\_CONST 0

→ LOAD\_CONST 1

BINARY\_ADD

POP\_TOP




1 + 21;

LOAD\_CONST 0

LOAD\_CONST 1

→ BINARY\_ADD

POP\_TOP



1 + 21;

LOAD\_CONST 0

LOAD\_CONST 1

BINARY\_ADD

→ POP\_TOP



```
class Statement(Node):
    def compile(self, ctx):
        self.expr.compile(ctx)
        ctx.emit(POP_TOP)

class Number(Node):
    def compile(self, ctx):
        ctx.emit(LOAD_CONST, ctx.new_const(JSNumber(self.value)))
```

```
class BinOp(Node):
    def compile(self, ctx):
        self.left.compile(ctx)
        self.right.compile(ctx)
        opname = {
            "+": "BINARY_ADD",
        }
        ctx.emit(opname[self.op])
```



# Object Model

```
class JSObject(object):  
    pass
```

```
class JSNumber(JSObject):  
    def __init__(self, value):  
        self.value = value
```

```
class JSNumber(JSObject):  
    def add(self, other):  
        assert isinstance(other, JSNumber)  
        return JSNumber(self.value + other.value)
```

# Interpreter (aka the Holy Grail)

```
class Interpreter(object):  
    def interpret(self):  
        pc = 0  
        while pc < len(self.bytecode):  
            opcode = self.bytecode[pc]  
            opname = OPCODE_TO_NAME[opcode]  
            pc = getattr(self, opname)(pc)
```

```
class Interpreter(object):  
    def LOAD_CONST(self, pc):  
        arg = ord(self.bytecode[pc + 1])  
        self.push(self.consts[arg])  
        return pc + 2
```

```
class Interpreter(object):  
    def BINARY_ADD(self, pc):  
        right = self.pop()  
        left = self.pop()  
        self.push(left.add(right))  
    return pc + 1
```

# Left as an exercise

- (PS: Come and find me or other PyPy/interpreter geeks to chat about these)
- if/while
- More types (strings!)
- More everything



# Useful references

- Topaz: <http://github.com/topazproject/topaz>
- Kermit Example Interpreter: <https://bitbucket.org/pypy/example-interpreter/>
- #pypy #topaz on Freenode

# Recap

- Lexer
- Parser
- AST
- Compiler
- Object Model
- Interpreter

Go forth and have fun

# Thanks, have an amazing PyCon!

<http://speakerdeck.com/u/alex>